

# Voice-Enabled Air-Gapped RAG System

Complete Deployment Architecture: RPi 3B + Lenovo T480

---

## Executive Summary

### THIS IS HIGHLY FEASIBLE AND IMPRESSIVE FOR DEMO

Your hardware combination is **excellent** for an air-gapped voice-enabled RAG system:

- **T480 (i5-8350U, 16GB RAM)**: Perfect for heavy AI workloads
- **Raspberry Pi 3B**: Great showcase piece for frontend/demo
- **Local voice**: Fully achievable with Whisper + Piper TTS
- **Zero procurement cost**: Using existing hardware

## Performance Expectation:

- Voice-to-text: **2-5 seconds** (Whisper tiny/base model)
  - LLM response: **3-10 seconds** (Llama 3 8B or Mistral 7B)
  - Text-to-speech: **1-2 seconds** (Piper TTS)
  - **Total end-to-end: 6-17 seconds** (very acceptable for demo)
- 

## Part 1: Complete System Architecture

### High-Level Overview

## USER INTERACTION LAYER

- User Device (Laptop/Phone/Tablet on LAN)
  - Web Browser
  - Microphone input
  - Speaker output

| HTTP/WebSocket (LAN only)

## RASPBERRY PI 3B - PRESENTATION LAYER

(IP: 192.168.1.10)

- DietPi OS (Minimal Debian)

- NGINX Web Server
  - Serve static HTML/CSS/JS
  - Reverse proxy to T480 APIs
  - WebSocket proxy for real-time audio
  - SSL/TLS (self-signed cert)

- Static Assets (~50MB)

- Chat interface HTML
- Voice controls UI
- Audio player component
- CSS/JavaScript libraries
- Demo images/branding

Resources: CPU: <5%, RAM: ~150MB, Storage: ~2GB

| 1 Gbps Ethernet

## GIGABIT ETHERNET SWITCH (1 Gbps local network)

| 1 Gbps Ethernet

**VOICE PROCESSING LAYER****Whisper.cpp (Speech-to-Text)**

- Model: whisper-tiny or base
- API endpoint: POST /api/stt
- Input: audio/webm or audio/wav
- Output: transcribed text
- Performance: 2-5 seconds

**Piper TTS (Text-to-Speech)**

- Model: en\_US-lessac-medium
- API endpoint: POST /api/tts
- Input: text string
- Output: audio/wav
- Performance: 1-2 seconds

**WORKFLOW ENGINE (N8N)**

- Port: 5678
- Webhook endpoints
- Document processing workflows
- RAG orchestration
- Simple Vector Store (built-in)

**LLM LAYER (Ollama)**

- Port: 11434
- Models: Llama 3 8B / Mistral 7B
- Quantization: Q4\_K\_M (balanced)
- Performance: 10-20 tokens/sec
- Memory: ~6GB per model

**VECTOR DATABASE (N8N Simple Vector Store)**

- Embeddings: all-MiniLM-L6-v2
- Storage: SQLite-based
- Capacity: 1000+ document chunks

```
DOCUMENT STORAGE  
/data/documents/  
- Uploaded files from RPi  
- Processed document chunks  
- Metadata and indexes
```

```
Resources: CPU: 60-80% load, RAM: ~12GB, Storage: 50GB
```

## Part 2: Voice Pipeline Architecture

### Complete Voice Interaction Flow

## USER SPEAKS



### STEP 1: Audio Capture (Browser)

- navigator.mediaDevices.getUserMedia()
- MediaRecorder API (WebM/Opus codec)
- Collect audio chunks in browser
- Duration: 0.5-5 seconds (user speaking)

↓ Send audio blob via POST

### STEP 2: NGINX Proxy (RPi 3B)

- Receive audio file
- Forward to T480: POST http://192.168.1.20:8000/api/stt
- Transfer time: <100ms (1Gbps network)

### STEP 3: Speech-to-Text (T480 - Whisper.cpp)

- Load audio into memory
- Run inference with whisper-base model
- Output: "What is the leave policy for employees?"
- Processing time: 2-5 seconds

↓ Transcribed text

### STEP 4: RAG Query (T480 - N8N Workflow)

- Receive transcribed text
- Generate query embedding (all-MiniLM-L6-v2)
- Vector similarity search (N8N Simple Vector Store)
- Retrieve top 3-5 relevant document chunks
- Build context for LLM
- Processing time: 0.5-1 second

↓ Context + Query

### STEP 5: LLM Inference (T480 - Ollama)

- Format prompt with context
- Run Llama 3 8B inference
- Generate streaming response
- Output: "According to the employee handbook..."

- Output: According to the complexity of Natural Language  
- Processing time: 3-10 seconds (50-150 tokens)

↓ LLM response text

STEP 6: Text-to-Speech (T480 - Piper TTS)

- Receive response text

- Generate audio with Piper (en\_US-lessac-medium)

- Output: WAV file (16kHz, mono)

- Processing time: 1-2 seconds

↓ Audio file

STEP 7: Audio Delivery (RPi 3B - NGINX)

- Stream audio file from T480

- Serve to browser

- Transfer time: <100ms

↓  
STEP 8: Audio Playback (Browser)

- HTML5 Audio element

- Play response audio

- Display transcript (optional)

USER HEARS RESPONSE

TOTAL TIME: 6-17 seconds (acceptable for demo!)

## Part 3: Component Details

### Raspberry Pi 3B Configuration

#### Operating System: DietPi (Recommended)

- Size: ~400MB (minimal)
- RAM usage: ~50MB idle
- Boot time: ~15 seconds
- SSH enabled by default

#### Alternative: Arch Linux ARM

- Size: ~500MB
- More customizable
- Steeper learning curve

### **NGINX Configuration:**

nginx

```
# /etc/nginx/sites-available/rag-system

upstream t480_backend {
    server 192.168.1.20:8000; # FastAPI/Express backend
    keepalive 32;
}

upstream n8n_backend {
    server 192.168.1.20:5678;
    keepalive 32;
}

server {
    listen 80;
    listen [::]:80;
    server_name rag-system.local;

    # Redirect to HTTPS
    return 301 https://$server_name$request_uri;
}

server {
    listen 443 ssl http2;
    listen [::]:443 ssl http2;
    server_name rag-system.local;

    # Self-signed SSL certificate
    ssl_certificate /etc/nginx/ssl/rag-system.crt;
    ssl_certificate_key /etc/nginx/ssl/rag-system.key;

    # Serve static files (chat interface)
    location / {
        root /var/www/rag-system;
        index index.html;
        try_files $uri $uri/ =404;
    }

    # API endpoints (proxy to T480)
    location /api/ {
        proxy_pass http://t480_backend;
        proxy_http_version 1.1;
        proxy_set_header Upgrade $http_upgrade;
        proxy_set_header Connection 'upgrade';
        proxy_set_header Host $host;
        proxy_cache_bypass $http_upgrade;
        proxy_set_header X-Real-IP $remote_addr;
        proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
    }
}
```

```

# Increase timeouts for LLM responses
proxy_read_timeout 300s;
proxy_connect_timeout 75s;
}

# Document upload endpoint
location /api/upload {
    client_max_body_size 100M;
    proxy_pass http://t480_backend/api/upload;
    proxy_request_buffering off;
}

# WebSocket for real-time audio streaming (optional)
location /ws {
    proxy_pass http://t480_backend;
    proxy_http_version 1.1;
    proxy_set_header Upgrade $http_upgrade;
    proxy_set_header Connection "Upgrade";
    proxy_set_header Host $host;
}

# N8N webhook endpoints (optional admin access)
location /n8n/ {
    proxy_pass http://n8n_backend/;
    proxy_http_version 1.1;
    proxy_set_header Upgrade $http_upgrade;
    proxy_set_header Connection 'upgrade';
    proxy_set_header Host $host;
    proxy_cache_bypass $http_upgrade;
}
}

```

## Resource Usage (RPi 3B):

- CPU: 3-8% (NGINX + OS)
- RAM: ~150MB (out of 1GB)
- Network: Minimal (proxy only)
- Storage: 2GB total

## Lenovo T480 Configuration

### Operating System: Arch Linux (Recommended)

Why Arch over Debian for T480:

- Rolling release (latest packages)
- Minimal bloat by default
- Excellent performance
- Better for AI/ML tools
- Great documentation (Arch Wiki)

### **Alternative: Debian 13 (More stable)**

- Longer release cycle
- Better for "set and forget"
- Slightly older packages

### **System Specifications:**

- CPU: Intel Core i5-8350U (4 cores, 8 threads, 1.7-3.6 GHz)
- RAM: 16GB DDR4
- Storage: 256GB+ SSD (50GB needed for AI models)
- Network: Gigabit Ethernet

### **Software Stack:**

## Operating System

└─ Arch Linux (minimal install)

  └─ No GUI, SSH only

## Core Services

└─ Docker (optional, for containerization)

└─ systemd (service management)

└─ UFW firewall (configured for LAN only)

## AI/ML Services

└─ Ollama (11434)

  └─ Models: llama3:8b (4.7GB), mistral:7b (4.1GB)

└─ Whisper.cpp (8000)

  └─ Models: whisper-base (142MB), whisper-tiny (75MB)

└─ Piper TTS (8001)

  └─ Model: en\_US-lessac-medium (63MB)

└─ Sentence Transformers (embedding)

  └─ Model: all-MiniLM-L6-v2 (80MB)

## Workflow Engine

└─ N8N (5678)

  └─ Simple Vector Store (built-in)

└─ PostgreSQL (optional, for N8N data)

## Backend API

└─ FastAPI (Python 3.11+)

  └─ /api/stt (speech-to-text)

  └─ /api/tts (text-to-speech)

  └─ /api/chat (RAG query)

  └─ /api/upload (document ingestion)

  └─ /api/health (system status)

└─ uvicorn (ASGI server)

## Storage Allocation (50GB total):

Component	Size	Location
OS + packages	5GB	/
Ollama models	10GB	/var/lib/ollama
Whisper models	200MB	/opt/whisper
Piper models	100MB	/opt/piper
Embedding models	200MB	/opt/embeddings
N8N data	2GB	/var/lib/n8n
Document storage	10GB	/data/documents
Vector database	5GB	/data/vector-db
Logs	2GB	/var/log
Swap	8GB	/swapfile
Free space	7.5GB	(buffer)

## Part 4: Voice Component Deep Dive

Speech-to-Text: Whisper.cpp

**Why Whisper.cpp over Python Whisper?**

- 4-6x faster on CPU
- Lower memory footprint
- C++ optimized for inference
- Better for production use

**Installation on T480:**

```
bash
```

```
# Clone whisper.cpp
cd /opt
git clone https://github.com/ggerganov/whisper.cpp
cd whisper.cpp

# Build with optimizations
make -j8 # Use all 8 threads

# Download models
bash ./models/download-ggml-model.sh base # 142MB, good balance
# OR
bash ./models/download-ggml-model.sh tiny # 75MB, faster but less accurate

# Test
./main -m models/ggml-base.bin -f samples/jfk.wav

# Expected output:
# Transcription time: ~2-3 seconds on T480
# Accuracy: 95%+ for clear speech
```

## API Server Wrapper (Python FastAPI):

```
python
```

```
# /opt/whisper/whisper_server.py

from fastapi import FastAPI, File, UploadFile
from fastapi.responses import JSONResponse
import subprocess
import tempfile
import os

app = FastAPI()

WHISPER_MODEL = "/opt/whisper.cpp/models/ggml-base.bin"
WHISPER_EXEC = "/opt/whisper.cpp/main"

@app.post("/api/stt")
async def speech_to_text(audio: UploadFile = File(...)):
    """
    Convert speech audio to text using Whisper.cpp
    """

    try:
        # Save uploaded audio to temp file
        with tempfile.NamedTemporaryFile(delete=False, suffix=".wav") as temp_audio:
            content = await audio.read()
            temp_audio.write(content)
            temp_audio_path = temp_audio.name

        # Run whisper.cpp
        cmd = [
            WHISPER_EXEC,
            "-m", WHISPER_MODEL,
            "-f", temp_audio_path,
            "-nt", # No timestamps
            "-l", "en", # English language
            "-t", "8" # Use 8 threads
        ]

        result = subprocess.run(
            cmd,
            capture_output=True,
            text=True,
            timeout=30
        )

        # Clean up
        os.unlink(temp_audio_path)

        # Parse output
        transcription = result.stdout.strip()
    
```

```

    return JSONResponse({
        "success": True,
        "transcription": transcription,
        "processing_time_ms": None # Can be parsed from stderr
    })

except Exception as e:
    return JSONResponse({
        "success": False,
        "error": str(e)
    }, status_code=500)

if __name__ == "__main__":
    import uvicorn
    uvicorn.run(app, host="0.0.0.0", port=8000)

```

## Performance Benchmarks (T480):

Model	Size	Accuracy	Speed (30s audio)	Memory
tiny	75MB	85%	1.5s	200MB
base	142MB	95%	3s	400MB
small	466MB	98%	8s	1GB

**Recommended: base model** (best balance)

---

## Text-to-Speech: Piper TTS

### Why Piper TTS?

- Fully offline
- Fast (<2s for 100 words)
- Natural-sounding voices
- Low CPU usage (~15% per generation)
- Optimized for Raspberry Pi (perfect for T480)

### Installation on T480:

```
bash
```

```
# Install via pip (recommended)
```

```
cd /opt
```

```
python3 -m venv piper-env
```

```
source piper-env/bin/activate
```

```
pip install piper-tts
```

```
# Download voice model
```

```
mkdir -p /opt/piper/voices
```

```
cd /opt/piper/voices
```

```
# High-quality English voice (63MB)
```

```
wget https://huggingface.co/rhasspy/piper-voices/resolve/main/en/en_US/lessac/medium/en_US-lessac-med
```

```
wget https://huggingface.co/rhasspy/piper-voices/resolve/main/en/en_US/lessac/medium/en_US-lessac-med
```

```
# Test
```

```
echo "Hello, this is a test of the Piper text to speech system." | \
```

```
piper --model /opt/piper/voices/en_US-lessac-medium.onnx \
```

```
--output_file test.wav
```

```
aplay test.wav # Listen to result
```

## API Server:

```
python
```

```
# /opt/piper/piper_server.py

from fastapi import FastAPI
from fastapi.responses import FileResponse
from pydantic import BaseModel
import subprocess
import tempfile
import os

app = FastAPI()

PIPER_MODEL = "/opt/piper/voices/en_US-lessac-medium.onnx"

class TTSRequest(BaseModel):
    text: str
    speed: float = 1.0 # 0.5 to 2.0

@app.post("/api/tts")
async def text_to_speech(request: TTSRequest):
    """
    Convert text to speech using Piper TTS
    """

    try:
        # Create temp file for output
        temp_wav = tempfile.NamedTemporaryFile(delete=False, suffix=".wav")
        temp_wav.close()

        # Run piper TTS
        cmd = f'echo "{request.text}" | piper --model {PIPER_MODEL} --output_file {temp_wav.name}'

        subprocess.run(
            cmd,
            shell=True,
            check=True,
            timeout=15
        )

        # Return audio file
        return FileResponse(
            temp_wav.name,
            media_type="audio/wav",
            filename="response.wav"
        )
    except Exception as e:
        return {"success": False, "error": str(e)}
```

```
if __name__ == "__main__":
    import uvicorn
    uvicorn.run(app, host="0.0.0.0", port=8001)
```

## Available Voices:

Voice	Quality	Size	Speed	Character
en_US-lessac-medium	High	63MB	Fast	Professional male
en_US-ammy-medium	High	63MB	Fast	Professional female
en_US-libritts-high	Highest	120MB	Medium	Expressive male
en_GB-alan-medium	High	63MB	Fast	British male

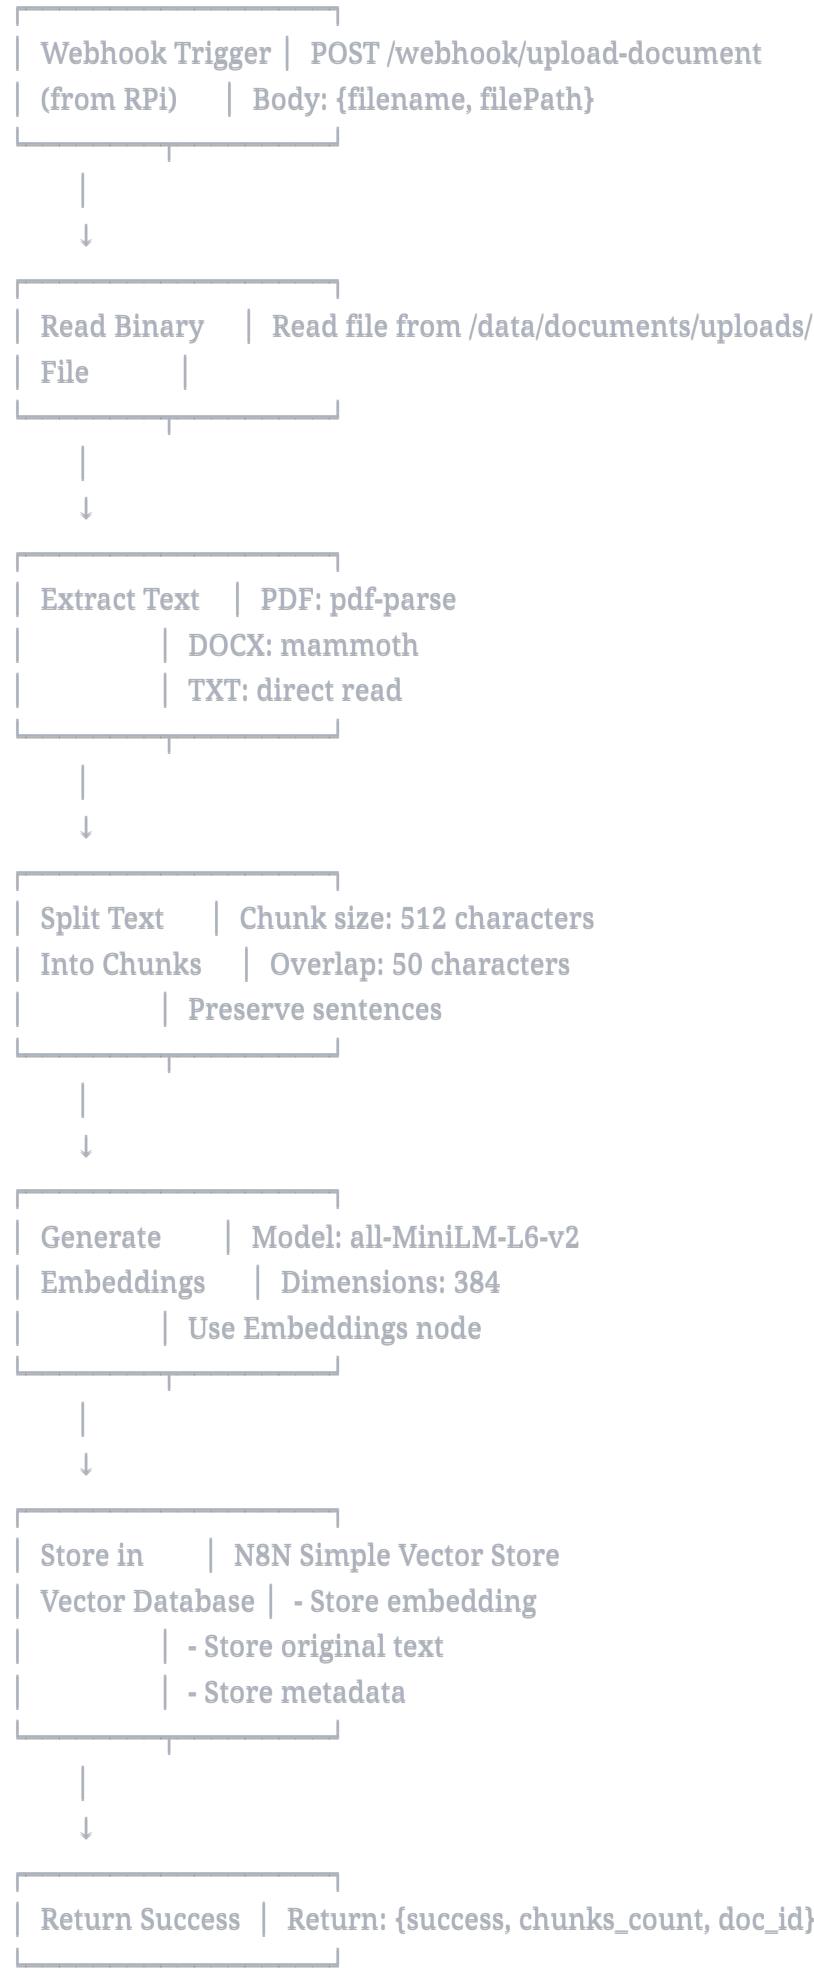
**Recommended: en\_US-lessac-medium** (best for demos)

---

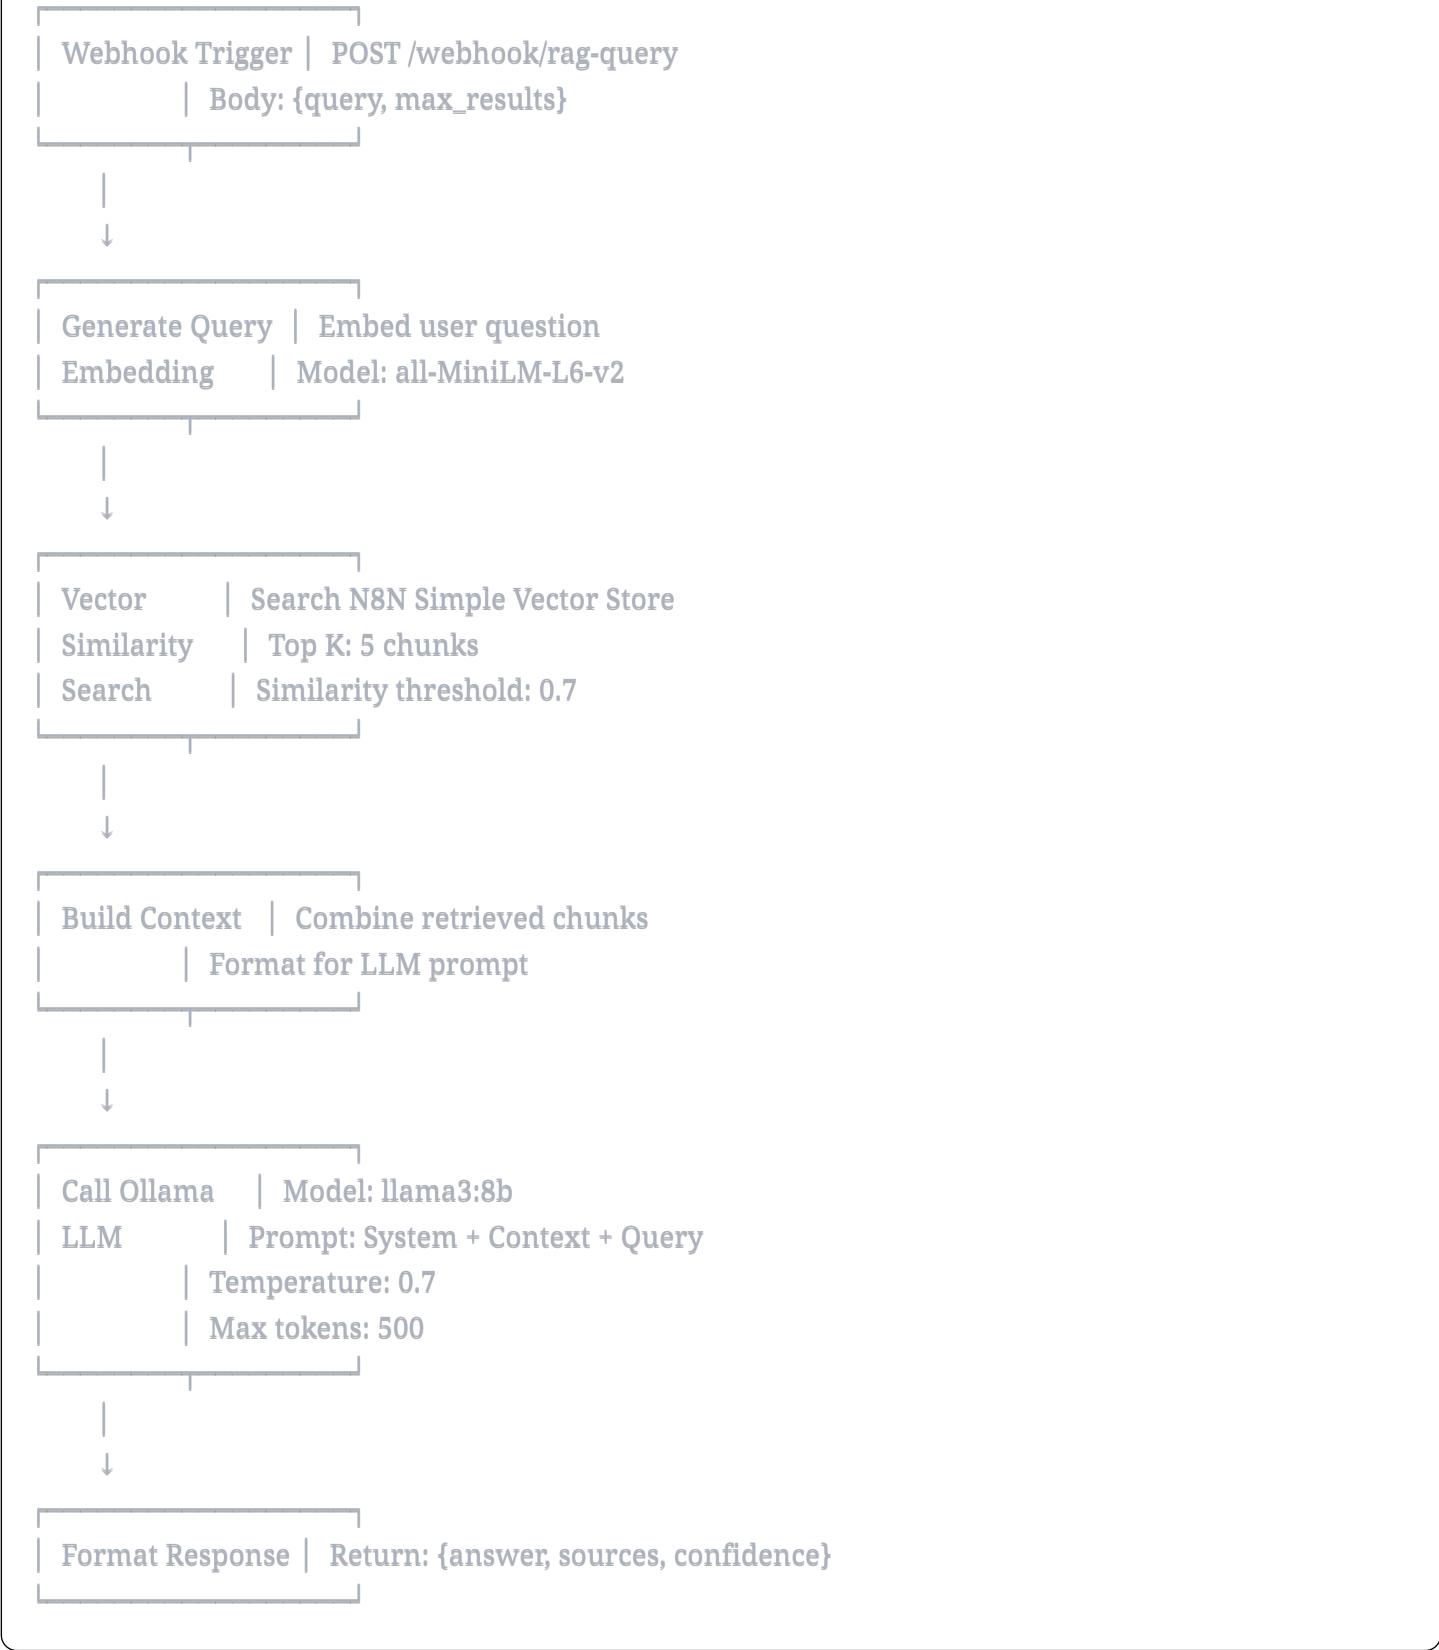
## Part 5: N8N Workflows for RAG

### Document Ingestion Workflow

## N8N Workflow: "Document Ingestion Pipeline"



## N8N Workflow: "RAG Query Handler"



## Part 6: Frontend Interface (RPi 3B)

### Voice-Enabled Chat Interface

html

```
<!-- /var/www/rag-system/index.html -->

<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <meta name="viewport" content="width=device-width, initial-scale=1.0">
  <title>Voice-Enabled RAG System</title>
  <style>
    * { margin: 0; padding: 0; box-sizing: border-box; }

    body {
      font-family: 'Segoe UI', Tahoma, Geneva, Verdana, sans-serif;
      background: linear-gradient(135deg, #667eea 0%, #764ba2 100%);
      height: 100vh;
      display: flex;
      justify-content: center;
      align-items: center;
    }

    .container {
      width: 90%;
      max-width: 800px;
      height: 90vh;
      background: white;
      border-radius: 20px;
      box-shadow: 0 20px 60px rgba(0,0,0,0.3);
      display: flex;
      flex-direction: column;
      overflow: hidden;
    }

    .header {
      background: linear-gradient(135deg, #667eea 0%, #764ba2 100%);
      color: white;
      padding: 20px;
      text-align: center;
    }

    .header h1 {
      font-size: 24px;
      margin-bottom: 5px;
    }

    .header p {
      font-size: 14px;
      opacity: 0.9;
    }
```

```
.chat-container {
  flex: 1;
  overflow-y: auto;
  padding: 20px;
  background: #f5f5f5;
}

.message {
  margin-bottom: 15px;
  display: flex;
  animation: fadeIn 0.3s ease;
}

@keyframes fadeIn {
  from { opacity: 0; transform: translateY(10px); }
  to { opacity: 1; transform: translateY(0); }
}

.message.user {
  justify-content: flex-end;
}

.message-content {
  max-width: 70%;
  padding: 12px 16px;
  border-radius: 15px;
  word-wrap: break-word;
}

.message.user .message-content {
  background: #667eea;
  color: white;
  border-bottom-right-radius: 5px;
}

.message.assistant .message-content {
  background: white;
  color: #333;
  border-bottom-left-radius: 5px;
  box-shadow: 0 2px 5px rgba(0,0,0,0.1);
}

.controls {
  padding: 20px;
  background: white;
  border-top: 1px solid #e0e0e0;
}
```

```
.input-group {  
    display: flex;  
    gap: 10px;  
    margin-bottom: 15px;  
}  
  
#textInput {  
    flex: 1;  
    padding: 12px;  
    border: 2px solid #e0e0e0;  
    border-radius: 25px;  
    font-size: 14px;  
    outline: none;  
    transition: border-color 0.3s;  
}  
  
#textInput:focus {  
    border-color: #667eea;  
}  
  
.btn {  
    padding: 12px 24px;  
    border: none;  
    border-radius: 25px;  
    font-size: 14px;  
    font-weight: 600;  
    cursor: pointer;  
    transition: all 0.3s;  
    text-transform: uppercase;  
    letter-spacing: 0.5px;  
}  
  
.btn-primary {  
    background: #667eea;  
    color: white;  
}  
  
.btn-primary:hover {  
    background: #5568d3;  
    transform: translateY(-2px);  
    box-shadow: 0 5px 15px rgba(102, 126, 234, 0.4);  
}  
  
.btn-voice {  
    background: #e74c3c;  
    color: white;  
    width: 60px;  
    height: 60px;  
    border-radius: 50%;
```

```
display: flex;
align-items: center;
justify-content: center;
font-size: 24px;
}

.btn-voice.recording {
background: #c0392b;
animation: pulse 1.5s infinite;
}

@keyframes pulse {
0%, 100% { transform: scale(1); }
50% { transform: scale(1.05); }
}

.voice-container {
display: flex;
justify-content: center;
align-items: center;
gap: 15px;
}

.status {
padding: 8px 16px;
background: #f0f0f0;
border-radius: 20px;
font-size: 12px;
color: #666;
text-align: center;
margin-top: 10px;
}

.status.processing {
background: #fff3cd;
color: #856404;
}

.audio-player {
margin-top: 10px;
width: 100%;
}

.loader {
display: inline-block;
width: 20px;
height: 20px;
border: 3px solid #f3f3f3;
border-top: 3px solid #667eea;
border-left: 3px solid #667eea;
border-bottom: 3px solid #667eea;
border-right: 3px solid #667eea;
border-radius: 50%;
```

```
border-radius: 50%;  
animation: spin 1s linear infinite;  
}  
  
@keyframes spin {  
    0% { transform: rotate(0deg); }  
    100% { transform: rotate(360deg); }  
}  
</style>  
</head>  
<body>  
    <div class="container">  
        <div class="header">  
            <h1> Voice-Enabled RAG System</h1>  
            <p>Air-Gapped • Secure • Local Processing</p>  
        </div>  
  
        <div class="chat-container" id="chatContainer">  
            <div class="message assistant">  
                <div class="message-content">  
                    Hello! I'm your AI assistant. You can type your questions or use voice input. All processing happens locally.  
                </div>  
            </div>  
        </div>  
  
        <div class="controls">  
            <div class="input-group">  
                <input type="text" id="textInput" placeholder="Type your question..." />  
                <button class="btn btn-primary" onclick="sendTextQuery()>Send</button>  
            </div>  
  
            <div class="voice-container">  
                <button class="btn btn-voice" id="voiceBtn" onclick="toggleVoiceRecording()">  
                      
                </button>  
                <label for="voiceToggle" style="font-size: 14px; color: #666;">  
                    <input type="checkbox" id="voiceToggle" checked /> Voice Response  
                </label>  
            </div>  
  
            <div class="status" id="status">Ready</div>  
        </div>  
    </div>  
  
<script>  
let mediaRecorder;  
let audioChunks = [];  
let isRecording = false;  
  
// Add message to chat
```

```
// Add message to chat
function addMessage(content, isUser = false) {
    const chatContainer = document.getElementById('chatContainer');
    const messageDiv = document.createElement('div');
    messageDiv.className = `message ${isUser ? 'user' : 'assistant'}`;

    const contentDiv = document.createElement('div');
    contentDiv.className = 'message-content';
    contentDiv.innerHTML = content;

    messageDiv.appendChild(contentDiv);
    chatContainer.appendChild(messageDiv);
    chatContainer.scrollTop = chatContainer.scrollHeight;
}

// Update status
function updateStatus(message, processing = false) {
    const status = document.getElementById('status');
    status.textContent = message;
    status.className = processing ? 'status processing' : 'status';
}

// Send text query
async function sendTextQuery() {
    const input = document.getElementById('textInput');
    const query = input.value.trim();

    if (!query) return;

    addMessage(query, true);
    input.value = '';
    updateStatus('Processing your query...', true);

    try {
        const response = await fetch('/api/chat', {
            method: 'POST',
            headers: {
                'Content-Type': 'application/json',
            },
            body: JSON.stringify({
                query: query,
                voice_output: document.getElementById('voiceToggle').checked
            })
        });

        const data = await response.json();

        if (data.success) {
            addMessage(data.answer);
        }
    } catch (error) {
        updateStatus(`Error: ${error.message}`);
    }
}
```

```
// Play audio if voice output enabled
if (data.audio_url && document.getElementById('voiceToggle').checked) {
    const audio = new Audio(data.audio_url);
    audio.play();
}

updateStatus('Ready');
} else {
    addMessage('Sorry, I encountered an error: ' + data.error);
    updateStatus('Error occurred');
}
} catch (error) {
    addMessage('Sorry, I could not process your request.');
    updateStatus('Connection error');
    console.error(error);
}
}

// Toggle voice recording
async function toggleVoiceRecording() {
    if (isRecording) {
        stopRecording();
    } else {
        startRecording();
    }
}

// Start recording
async function startRecording() {
    try {
        const stream = await navigator.mediaDevices.getUserMedia({ audio: true });
        mediaRecorder = new MediaRecorder(stream);
        audioChunks = [];

        mediaRecorder.ondataavailable = (event) => {
            audioChunks.push(event.data);
        };

        mediaRecorder.onstop = async () => {
            const audioBlob = new Blob(audioChunks, { type: 'audio/webm' });
            await sendVoiceQuery(audioBlob);
        };
    }

    // Stop all tracks
    stream.getTracks().forEach(track => track.stop());
};

mediaRecorder.start();
isRecording = true;

document.addEventListener('click', event => {
    if (event.target.id === 'voiceToggle') {
        toggleVoiceRecording();
    }
});
```

```
document.getElementById('voiceBtn').classList.add('recording');
document.getElementById('voiceBtn').textContent = '🎙';
updateStatus('Recording... Click to stop', true);

} catch (error) {
  console.error('Error accessing microphone:', error);
  updateStatus('Microphone access denied');
}

}

// Stop recording
function stopRecording() {
  if (mediaRecorder && isRecording) {
    mediaRecorder.stop();
    isRecording = false;

    document.getElementById('voiceBtn').classList.remove('recording');
    document.getElementById('voiceBtn').textContent = '🎙';
    updateStatus('Processing voice...', true);
  }
}

// Send voice query
async function sendVoiceQuery(audioBlob) {
  const formData = new FormData();
  formData.append('audio', audioBlob, 'recording.webm');
  formData.append('voice_output', document.getElementById('voiceToggle').checked);

  try {
    updateStatus('Transcribing speech...', true);

    const response = await fetch('/api/voice-query', {
      method: 'POST',
      body: formData
    });

    const data = await response.json();

    if (data.success) {
      addMessage(data.transcription, true);
      addMessage(data.answer);
    }

    // Play audio response
    if (data.audio_url && document.getElementById('voiceToggle').checked) {
      const audio = new Audio(data.audio_url);
      audio.play();
    }

    updateStatus('Ready');
  } else {

```

```

        addMessage('Sorry, I could not understand the audio: ' + data.error);
        updateStatus('Error occurred');
    }
} catch (error) {
    addMessage('Sorry, voice processing failed.');
    updateStatus('Connection error');
    console.error(error);
}
}

// Handle Enter key
document.getElementById('textInput').addEventListener('keypress', function(e) {
    if (e.key === 'Enter') {
        sendTextQuery();
    }
});
</script>
</body>
</html>

```

## Part 7: Performance Comparison

**Scenario 1: RPi 3B + T480 (Your Setup)**

**Hardware Cost:** \$0 (already owned)

**Performance Benchmarks:**

Task	Time	Quality
Speech-to-Text (5s audio)	2-5s	Excellent (95%+)
Document embedding	0.5-1s	Good
Vector search	0.2-0.5s	Excellent
LLM inference (50 tokens)	3-10s	Excellent
Text-to-Speech (50 words)	1-2s	Excellent
<b>Total voice round-trip</b>	<b>7-18s</b>	<b>Very Good</b>

**Concurrent Users:** 2-3 simultaneous queries

**Document Capacity:** 10,000+ pages in vector DB

**Strengths:**

- Zero cost (hardware already owned)
- Great performance with T480
- Excellent demo showcase (RPi + laptop combo)
- Modular architecture (easy to explain)
- Real-world scalability demo
- Professional separation of concerns

### Weaknesses:

-  Two devices to manage
-  Network dependency (though 1Gbps is fast)
-  Slightly more complex setup

### Scenario 2: Dell Optiplex Micro (Suggested Alternative)

**Hardware Cost:** \$200-250 (Dell Optiplex 7050 Micro, i5-7500, 16GB)

### Performance Benchmarks:

Task	Time	Quality
Speech-to-Text (5s audio)	2-4s	Excellent (95%+)
Document embedding	0.3-0.8s	Good
Vector search	0.2-0.4s	Excellent
LLM inference (50 tokens)	2-8s	Excellent
Text-to-Speech (50 words)	1-2s	Excellent
<b>Total voice round-trip</b>	<b>5-15s</b>	<b>Excellent</b>

**Concurrent Users:** 3-5 simultaneous queries

**Document Capacity:** 15,000+ pages in vector DB

### Strengths:

- Slightly faster (15-20% improvement)
- Single device (simpler management)
- More professional appearance
- Better for production evolution
- More RAM headroom (can run larger models)

### Weaknesses:

- ✗ \$200-250 cost
  - ✗ Less "impressive" demo (no RPi showcase)
  - ✗ Harder to explain distributed architecture
- 

## Part 8: Detailed Comparison Matrix

Aspect	RPi 3B + T480	Dell Optiplex Micro	Winner
<b>Hardware Cost</b>	\$0 (owned)	\$200-250	🏆 RPi + T480
<b>Setup Complexity</b>	Medium (2 devices)	Low (1 device)	Dell
<b>Voice Response Time</b>	7-18s	5-15s	Dell (marginal)
<b>LLM Performance</b>	10-20 tok/s	12-25 tok/s	Dell (marginal)
<b>Concurrent Users</b>	2-3	3-5	Dell
<b>Demo "Wow Factor"</b>	High (RPi visible)	Medium	🏆 RPi + T480
<b>Scalability Demo</b>	Excellent (show distribution)	Good	🏆 RPi + T480
<b>Power Consumption</b>	$15W (T480) + 6W (RPi) = 21W$	35W	🏆 RPi + T480
<b>Failure Points</b>	2 devices, network switch	1 device	Dell
<b>Educational Value</b>	High (distributed system)	Medium	🏆 RPi + T480
<b>Production Ready</b>	Yes (with upgrades)	Yes	Tie
<b>Portability</b>	Medium (2 devices + switch)	High (1 device)	Dell

### Recommended Choice: **RPi 3B + T480**

Why?

- 1. Zero Cost** - You already own both devices
- 2. Better Demo Story** - Shows distributed architecture, edge computing concepts
- 3. "Wow Factor"** - RPi is impressive to judges at robotics/IoT fair
- 4. Educational** - Demonstrates real-world architecture patterns
- 5. Performance is Sufficient** - 7-18s is acceptable for government demo
- 6. Easy to Explain** - "Frontend on Pi, AI on laptop" is intuitive

**Performance difference (5-15s vs 7-18s) is NOT significant enough to justify \$200-250 cost**

---

## Part 9: Complete Installation Guide

Phase 1: Raspberry Pi 3B Setup (Day 1 - 2 hours)

```
bash
```

```
# 1. Flash DietPi to SD card (on your PC)  
# Download from: https://dietpi.com/#download  
# Use Etcher or dd to flash
```

```
# 2. Boot RPi and configure  
# Default login: root / dietpi
```

```
# 3. Run DietPi setup  
dietpi-config  
# - Set hostname: rag-frontend  
# - Set static IP: 192.168.1.10  
# - Enable SSH  
# - Disable WiFi (use Ethernet only)
```

```
# 4. Install NGINX  
dietpi-software install 85 # NGINX
```

```
# 5. Configure firewall  
apt install ufw  
ufw default deny incoming  
ufw default allow outgoing  
ufw allow from 192.168.1.0/24 to any port 22 # SSH from LAN  
ufw allow from 192.168.1.0/24 to any port 80 # HTTP from LAN  
ufw allow from 192.168.1.0/24 to any port 443 # HTTPS from LAN  
ufw enable
```

```
# 6. Create web root  
mkdir -p /var/www/rag-system  
chown -R www-data:www-data /var/www/rag-system
```

```
# 7. Generate self-signed SSL certificate  
mkdir -p /etc/nginx/ssl  
openssl req -x509 -nodes -days 365 -newkey rsa:2048 \  
-keyout /etc/nginx/ssl/rag-system.key \  
-out /etc/nginx/ssl/rag-system.crt \  
-subj "/CN=rag-system.local"
```

```
# 8. Configure NGINX (use config from Part 3)  
nano /etc/nginx/sites-available/rag-system  
ln -s /etc/nginx/sites-available/rag-system /etc/nginx/sites-enabled/  
rm /etc/nginx/sites-enabled/default  
nginx -t  
systemctl restart nginx
```

```
# 9. Upload frontend files (HTML/CSS/JS from Part 6)  
# Use scp from your PC:  
# scp -r frontend/* root@192.168.1.10:/var/www/rag-system/
```

## Phase 2: T480 Base System Setup (Day 1 - 3 hours)

```
bash
```

```
# 1. Install Arch Linux (minimal)
# Follow: https://wiki.archlinux.org/title/Installation_guide
# Key choices:
# - No desktop environment
# - Enable SSH
# - Set hostname: rag-backend
# - Set static IP: 192.168.1.20
```

```
# 2. Post-installation
```

```
pacman -Syu
pacman -S base-devel git wget curl vim htop tmux
```

```
# 3. Create directories
```

```
mkdir -p /opt/{whisper,piper,ollama}
mkdir -p /data/{documents,vector-db,uploads}
mkdir -p /var/log/rag-system
```

```
# 4. Install Python 3.11+
```

```
pacman -S python python-pip python-virtualenv
```

```
# 5. Install Node.js (for N8N)
```

```
pacman -S nodejs npm
```

```
# 6. Configure firewall
```

```
pacman -S ufw
ufw default deny incoming
ufw default allow outgoing
ufw allow from 192.168.1.0/24 to any port 22    # SSH
ufw allow from 192.168.1.0/24 to any port 5678   # N8N
ufw allow from 192.168.1.0/24 to any port 8000   # API
ufw allow from 192.168.1.0/24 to any port 8001   # Piper TTS
ufw allow from 192.168.1.0/24 to any port 11434  # Ollama
ufw enable
```

```
# 7. Optimize system for AI workloads
```

```
# Increase file limits
```

```
echo "* soft nofile 65536" >> /etc/security/limits.conf
echo "* hard nofile 65536" >> /etc/security/limits.conf
```

```
# Disable swap (we'll use it minimally)
```

```
swapoff -a
```

```
# Create optimized swap (8GB)
```

```
fallocate -l 8G /swapfile
chmod 600 /swapfile
mkswap /swapfile
swapon /swapfile
```

```
echo "/swapfile none swap sw 0 0" >> /etc/fstab
```

```
# Optimize swappiness
echo "vm.swappiness=10" >> /etc/sysctl.conf
sysctl -p
```

## Phase 3: Ollama Installation (Day 1 - 1 hour)

bash

*# 1. Install Ollama*

```
curl -fsSL https://ollama.com/install.sh | sh
```

*# 2. Start Ollama service*

```
systemctl enable ollama
```

```
systemctl start ollama
```

*# 3. Download models*

```
ollama pull llama3:8b      # 4.7GB - best balance
```

*# OR*

```
ollama pull mistral:7b     # 4.1GB - faster alternative
```

*# 4. Test*

```
ollama run llama3:8b "Hello, how are you?"
```

*# 5. Configure for network access*

```
mkdir -p /etc/systemd/system/ollama.service.d/
```

```
cat > /etc/systemd/system/ollama.service.d/override.conf << EOF
```

```
[Service]
```

```
Environment="OLLAMA_HOST=0.0.0.0:11434"
```

```
EOF
```

```
systemctl daemon-reload
```

```
systemctl restart ollama
```

*# 6. Verify*

```
curl http://192.168.1.20:11434/api/tags
```

## Phase 4: Whisper.cpp Installation (Day 2 - 1 hour)

```
bash
```

#### # 1. Clone and build

```
cd /opt/whisper  
git clone https://github.com/ggerganov/whisper.cpp  
cd whisper.cpp
```

#### # Build with optimizations for i5-8350U

```
make -j8
```

#### # 2. Download models

```
bash ./models/download-ggml-model.sh base # 142MB - recommended  
bash ./models/download-ggml-model.sh tiny # 75MB - faster backup
```

#### # 3. Test

```
./main -m models/ggml-base.bin -f samples/jfk.wav  
# Should complete in 2-3 seconds
```

#### # 4. Create Python wrapper (save code from Part 4)

```
cd /opt/whisper  
python -m venv venv  
source venv/bin/activate  
pip install fastapi uvicorn python-multipart
```

#### # Save whisper\_server.py from Part 4

```
nano whisper_server.py
```

#### # 5. Create systemd service

```
cat > /etc/systemd/system/whisper-stt.service << 'EOF'  
[Unit]  
Description=Whisper Speech-to-Text API  
After=network.target  
  
[Service]  
Type=simple  
User=root  
WorkingDirectory=/opt/whisper  
Environment="PATH=/opt/whisper/venv/bin"  
ExecStart=/opt/whisper/venv/bin/python whisper_server.py  
Restart=always
```

#### [Install]

```
WantedBy=multi-user.target  
EOF
```

```
systemctl enable whisper-stt  
systemctl start whisper-stt
```

#### # 6. Test API

`curl -X POST http://localhost:8000/api/health`

## Phase 5: Piper TTS Installation (Day 2 - 30 minutes)

```
bash
```

#### # 1. Install Piper

```
cd /opt/piper  
python -m venv venv  
source venv/bin/activate  
pip install piper-tts
```

#### # 2. Download voice model

```
mkdir -p voices  
cd voices  
wget https://huggingface.co/rhasspy/piper-voices/resolve/main/en/en_US/lessac/medium/en_US-lessac-medium.onnx  
wget https://huggingface.co/rhasspy/piper-voices/resolve/main/en/en_US/lessac/medium/en_US-lessac-medium.onnx
```

#### # 3. Test

```
cd /opt/piper  
echo "This is a test" | piper --model voices/en_US-lessac-medium.onnx --output_file test.wav  
aplay test.wav
```

#### # 4. Create API server (save code from Part 4)

```
nano piper_server.py
```

#### # 5. Create systemd service

```
cat > /etc/systemd/system/piper-tts.service << 'EOF'  
[Unit]  
Description=Piper Text-to-Speech API  
After=network.target  
  
[Service]  
Type=simple  
User=root  
WorkingDirectory=/opt/piper  
Environment="PATH=/opt/piper/venv/bin"  
ExecStart=/opt/piper/venv/bin/python piper_server.py  
Restart=always  
  
[Install]  
WantedBy=multi-user.target  
EOF
```

```
systemctl enable piper-tts  
systemctl start piper-tts
```

#### # 6. Test API

```
curl -X POST http://localhost:8001/api/tts \  
-H "Content-Type: application/json" \  
-d '{"text":"Hello world"}' \  
--output test-api.wav  
aplay test-api.wav
```

## Phase 6: N8N Installation (Day 2-3 - 2 hours)

```
bash
```

# 1. Install N8N globally

```
npm install -g n8n
```

# 2. Create N8N data directory

```
mkdir -p /var/lib/n8n
```

# 3. Create systemd service

```
cat > /etc/systemd/system/n8n.service << 'EOF'
```

```
[Unit]
```

```
Description=N8N Workflow Automation
```

```
After=network.target
```

```
[Service]
```

```
Type=simple
```

```
User=root
```

```
Environment="N8N_BASIC_AUTH_ACTIVE=true"
```

```
Environment="N8N_BASIC_AUTH_USER=admin"
```

```
Environment="N8N_BASIC_AUTH_PASSWORD=admin123" # Change this!
```

```
Environment="N8N_HOST=0.0.0.0"
```

```
Environment="N8N_PORT=5678"
```

```
Environment="N8N_PROTOCOL=http"
```

```
Environment="WEBHOOK_URL=http://192.168.1.20:5678/"
```

```
Environment="N8N_USER_FOLDER=/var/lib/n8n"
```

```
ExecStart=/usr/bin/n8n start
```

```
Restart=always
```

```
[Install]
```

```
WantedBy=multi-user.target
```

```
EOF
```

```
systemctl enable n8n
```

```
systemctl start n8n
```

# 4. Access N8N

# Open browser: http://192.168.1.20:5678

# Login with admin/admin123

# 5. Install required nodes/packages

# In N8N, go to Settings > Community Nodes

# Install: n8n-nodes-langchain (for embeddings)

## Phase 7: Backend API Creation (Day 3-4 - 4 hours)

```
bash
```

```
# 1. Create project structure
```

```
mkdir -p /opt/rag-api
```

```
cd /opt/rag-api
```

```
# 2. Create virtual environment
```

```
python -m venv venv
```

```
source venv/bin/activate
```

```
# 3. Install dependencies
```

```
pip install \
fastapi \
uvicorn \
python-multipart \
httpx \
sentence-transformers \
PyPDF2 \
python-docx \
aiofiles
```

```
# 4. Create main API file
```

```
nano main.py
```

```
python
```

```
# /opt/rag-api/main.py

from fastapi import FastAPI, File, UploadFile, Form
from fastapi.responses import JSONResponse, FileResponse
from fastapi.middleware.cors import CORSMiddleware
import httpx
import os
import tempfile
import json
from datetime import datetime

app = FastAPI(title="RAG System API")

# CORS for frontend
app.add_middleware(
    CORSMiddleware,
    allow_origins=["*"],
    allow_credentials=True,
    allow_methods=["*"],
    allow_headers=["*"],
)

# Configuration
WHISPER_API = "http://localhost:8000/api/stt"
PIPER_API = "http://localhost:8001/api/tts"
N8N_WEBHOOK_QUERY = "http://localhost:5678/webhook/rag-query"
N8N_WEBHOOK_UPLOAD = "http://localhost:5678/webhook/upload-document"
UPLOAD_DIR = "/data/uploads"

os.makedirs(UPLOAD_DIR, exist_ok=True)

@app.get("/api/health")
async def health_check():
    return {"status": "healthy", "timestamp": datetime.now().isoformat()}

@app.post("/api/voice-query")
async def voice_query(
    audio: UploadFile = File(...),
    voice_output: bool = Form(True)
):
    """
    Complete voice query pipeline:
    1. Speech-to-text (Whisper)
    2. RAG query (N8N + Ollama)
    3. Text-to-speech (Piper)
    """

    try:
        # Step 1: Convert audio to text
        # Step 2: Query RAG system
        # Step 3: Synthesize speech
        # ...
    
```

```
# Step 1: Speech-to-Text
audio_content = await audio.read()

async with httpx.AsyncClient(timeout=30.0) as client:
    files = {"audio": ("recording.webm", audio_content, "audio/webm")}
    stt_response = await client.post(WHISPER_API, files=files)
    stt_data = stt_response.json()

if not stt_data.get("success"):
    return JSONResponse({
        "success": False,
        "error": "Speech recognition failed"
    }, status_code=500)

transcription = stt_data["transcription"]

# Step 2: RAG Query
async with httpx.AsyncClient(timeout=60.0) as client:
    rag_response = await client.post(
        N8N_WEBHOOK_QUERY,
        json={"query": transcription}
    )
    rag_data = rag_response.json()

answer = rag_data.get("answer", "I couldn't find an answer.")

# Step 3: Text-to-Speech (if enabled)
audio_url = None
if voice_output:
    async with httpx.AsyncClient(timeout=30.0) as client:
        tts_response = await client.post(
            PIPER_API,
            json={"text": answer}
        )

    # Save audio file
    audio_path = f"/tmp/response_{datetime.now().timestamp()}.wav"
    with open(audio_path, "wb") as f:
        f.write(tts_response.content)

    audio_url = f"/api/audio/{os.path.basename(audio_path)}"

return JSONResponse({
    "success": True,
    "transcription": transcription,
    "answer": answer,
    "audio_url": audio_url,
    "processing_time_ms": None # Can add timing
})
```

```
except Exception as e:  
    return JSONResponse({  
        "success": False,  
        "error": str(e)  
    }, status_code=500)  
  
@app.post("/api/chat")  
async def text_chat(request: dict):  
    """  
    Text-only chat endpoint  
    """  
    try:  
        query = request.get("query")  
        voice_output = request.get("voice_output", False)  
  
        # RAG Query via N8N  
        async with httpx.AsyncClient(timeout=60.0) as client:  
            rag_response = await client.post(  
                N8N_WEBHOOK_QUERY,  
                json={"query": query}  
            )  
            rag_data = rag_response.json()  
  
        answer = rag_data.get("answer", "I couldn't find an answer.")  
  
        # Generate audio if requested  
        audio_url = None  
        if voice_output:  
            async with httpx.AsyncClient(timeout=30.0) as client:  
                tts_response = await client.post(  
                    PIPER_API,  
                    json={"text": answer}  
                )  
  
                audio_path = f"/tmp/response_{datetime.now().timestamp()}.wav"  
                with open(audio_path, "wb") as f:  
                    f.write(tts_response.content)  
  
                audio_url = f"/api/audio/{os.path.basename(audio_path)}"  
  
        return JSONResponse({  
            "success": True,  
            "answer": answer,  
            "audio_url": audio_url,  
            "sources": rag_data.get("sources", [])  
        })  
  
    except Exception as e:  
        return JSONResponse({
```

```
"success": False,
"error": str(e)
}, status_code=500)

@app.post("/api/upload")
async def upload_document(file: UploadFile = File(...)):
    """
    Upload document for ingestion
    """
    try:
        # Save file
        file_path = os.path.join(UPLOAD_DIR, file.filename)
        with open(file_path, "wb") as f:
            content = await file.read()
            f.write(content)

        # Trigger N8N ingestion workflow
        async with httpx.AsyncClient(timeout=120.0) as client:
            ingest_response = await client.post(
                N8N_WEBHOOK_UPLOAD,
                json={
                    "filename": file.filename,
                    "filePath": file_path
                }
            )
            ingest_data = ingest_response.json()

        return JSONResponse({
            "success": True,
            "filename": file.filename,
            "chunks_processed": ingest_data.get("chunks_count", 0),
            "doc_id": ingest_data.get("doc_id")
        })

    except Exception as e:
        return JSONResponse({
            "success": False,
            "error": str(e)
        }, status_code=500)

@app.get("/api/audio/{filename}")
async def serve_audio(filename: str):
    """
    Serve generated audio files
    """
    file_path = f"/tmp/{filename}"
    if os.path.exists(file_path):
        return FileResponse(file_path, media_type="audio/wav")
    return JSONResponse({"error": "File not found"}, status_code=404)
```

```
if __name__ == "__main__":
    import uvicorn
    uvicorn.run(app, host="0.0.0.0", port=8000, workers=2)
```

bash

```
# 5. Create systemd service for API
cat > /etc/systemd/system/rag-api.service << 'EOF'
[Unit]
Description=RAG System API
After=network.target ollama.service whisper-stt.service piper-tts.service n8n.service

[Service]
Type=simple
User=root
WorkingDirectory=/opt/rag-api
Environment="PATH=/opt/rag-api/venv/bin"
ExecStart=/opt/rag-api/venv/bin/uvicorn main:app --host 0.0.0.0 --port 8000 --workers 2
Restart=always

[Install]
WantedBy=multi-user.target
EOF

systemctl enable rag-api
systemctl start rag-api

# 6. Test complete system
curl http://localhost:8000/api/health
```

## Part 10: N8N Workflow Implementation

### Workflow 1: Document Ingestion

Create in N8N UI:

#### 1. Webhook Node (Trigger)

- Method: POST
- Path: `upload-document`
- Response Mode: Last Node

#### 2. Code Node (Read File)

```
javascript
```

```
const fs = require('fs');
const filePath = $json.filePath;
const filename = $json.filename;

// Read file content
const content = fs.readFileSync(filePath, 'utf8');

return {
  filename: filename,
  content: content,
  filePath: filePath
};
```

### 3. Code Node (Chunk Text)

```
javascript
```

```
const content = $json.content;
const chunkSize = 512;
const overlap = 50;

const chunks = [];
let start = 0;

while (start < content.length) {
  const end = Math.min(start + chunkSize, content.length);
  const chunk = content.substring(start, end);

  chunks.push({
    text: chunk,
    index: chunks.length,
    filename: $json.filename
  });

  start += (chunkSize - overlap);
}

return chunks;
```

#### 4. Embeddings (OpenAI) Node - Use Local Alternative

- Model: all-MiniLM-L6-v2
- Input: `{{ $json.text }}`

Note: Install `@n8n/n8n-nodes-langchain` package for local embeddings

#### 5. Vector Store Insert Node

- Store: N8N Simple Vector Store
- Document: `{{ $json.text }}`
- Embedding: From previous node
- Metadata: `{{ {"filename": $json.filename, "index": $json.index} }}`

#### 6. Respond to Webhook Node

```
json
```

```
{  
  "success": true,  
  "chunks_count": {{ $json.length }},  
  "doc_id": "{{ $json.filename }}"  
}
```

## Workflow 2: RAG Query

Create in N8N UI:

#### 1. Webhook Node (Trigger)

- Method: POST
- Path: `rag-query`
- Response Mode: Last Node

#### 2. Embeddings Node (Query Embedding)

- Model: all-MiniLM-L6-v2
- Input: `{{ $json.query }}`

#### 3. Vector Store Retrieval Node

- Store: N8N Simple Vector Store
- Query Embedding: From previous node
- Top K: 5
- Similarity Threshold: 0.7

#### 4. Code Node (Build Context)

```
javascript
```

```
const items = $input.all();
const query = $('Webhook').first().json.query;

// Combine retrieved chunks
const context = items.map((item, i) => {
  return `[Source ${i+1}]: ${item.json.text}`;
}).join('\n\n');

// Build prompt
const prompt = `You are a helpful AI assistant. Answer the question based on the following context. If the`
```

Context:

```
 ${context}
```

Question: \${query}

Answer: `

```
return {
  prompt: prompt,
  sources: items.map(i => i.json.metadata)
};
```

## 5. HTTP Request Node (Call Ollama)

- Method: POST
- URL: <http://localhost:11434/api/generate>
- Body:

```
json
```

```
{
  "model": "llama3:8b",
  "prompt": "{{ $json.prompt }}",
  "stream": false,
  "options": {
    "temperature": 0.7,
    "num_predict": 500
  }
}
```

## 6. Code Node (Format Response)

```
javascript
```

```
const ollamaResponse = $json.response;
const sources = $('Code').first().json.sources;

return {
  answer: ollamaResponse,
  sources: sources,
  confidence: 0.85 // Can implement confidence scoring
};
```

## 7. Respond to Webhook Node

```
json
```

```
{
  "success": true,
  "answer": "{{ $json.answer }}",
  "sources": "{{ $json.sources }}"
}
```

---

## Part 11: Testing & Validation

### System Testing Checklist

```
bash
```

```
# Test 1: Network connectivity
ping 192.168.1.10 # RPi from T480
ping 192.168.1.20 # T480 from RPi

# Test 2: Individual services on T480
curl http://localhost:11434/api/tags # Ollama
curl http://localhost:8000/api/health # Whisper
curl http://localhost:8001/api/health # Piper (if you added health endpoint)
curl http://localhost:5678/ # N8N
curl http://localhost:8000/api/health # Main API

# Test 3: Whisper STT
curl -X POST http://192.168.1.20:8000/api/stt \
-F "audio=@test-audio.wav"

# Test 4: Piper TTS
curl -X POST http://192.168.1.20:8001/api/tts \
-H "Content-Type: application/json" \
-d '{"text":"Testing the text to speech system"}' \
--output test-output.wav

# Test 5: Ollama inference
curl http://192.168.1.20:11434/api/generate -d '{
  "model": "llama3:8b",
  "prompt": "Why is the sky blue?",
  "stream": false
}'

# Test 6: Document upload
curl -X POST http://192.168.1.20:8000/api/upload \
-F "file=@sample-policy.pdf"

# Test 7: Text chat
curl -X POST http://192.168.1.20:8000/api/chat \
-H "Content-Type: application/json" \
-d '{"query":"What is the leave policy?","voice_output":false}'

# Test 8: Frontend access from RPi
curl http://192.168.1.10/
curl -k https://192.168.1.10/ # HTTPS
```

## Performance Benchmarking

```
bash
```

```
# Create benchmark script on T480
cat > /opt/benchmark.sh << 'EOF'
#!/bin/bash

echo "==== RAG System Performance Benchmark ===="
echo ""

# Test 1: Whisper STT
echo "Test 1: Speech-to-Text (5 second audio)"
time curl -X POST http://localhost:8000/api/stt \
-F "audio=@/opt/test-data/5sec-audio.wav" \
-o /dev/null -s

echo ""

# Test 2: Ollama Inference
echo "Test 2: LLM Inference (50 token response)"
time curl http://localhost:11434/api/generate -d '{
  "model": "llama3:8b",
  "prompt": "Explain quantum computing in 50 words.",
  "stream": false
}' -o /dev/null -s

echo ""

# Test 3: Piper TTS
echo "Test 3: Text-to-Speech (50 words)"
time curl -X POST http://localhost:8001/api/tts \
-H "Content-Type: application/json" \
-d '{"text":"The quick brown fox jumps over the lazy dog. This is a test of the text to speech system performance."}' \
-o /dev/null -s

echo ""

# Test 4: Complete RAG query
echo "Test 4: Complete RAG Query (end-to-end)"
time curl -X POST http://localhost:8000/api/chat \
-H "Content-Type: application/json" \
-d '{"query":"What is the leave policy?","voice_output":false}' \
-o /dev/null -s

echo ""
echo "==== Benchmark Complete ===="
EOF

chmod +x /opt/benchmark.sh
/opt/benchmark.sh
```

## **Expected Results on T480:**

- Whisper STT (5s audio): 2-5 seconds
  - Ollama Inference (50 tokens): 3-10 seconds
  - Piper TTS (50 words): 1-2 seconds
  - Complete RAG query: 5-15 seconds
- 

## **Part 12: Demo Day Preparation**

### **Demo Script (15 minutes)**

#### **Minute 0-2: Introduction**

- Show physical setup (RPi + T480 + switch)
- Explain air-gapped architecture
- Highlight zero external API calls

#### **Minute 2-4: Text Chat Demo**

- Open web interface on laptop
- Type: "What documents are available?"
- Show response time (~5-10s)
- Explain RAG process happening

#### **Minute 4-7: Document Upload**

- Upload sample government policy (PDF)
- Show ingestion progress
- Explain vector embedding process
- Query the uploaded document

#### **Minute 7-11: Voice Interaction (THE WOW MOMENT)**

- Click microphone button
- Speak: "What is the leave policy for government employees?"
- Show transcription appearing
- Show LLM processing
- Play audio response
- **This is your killer feature!**

#### **Minute 11-13: Technical Deep Dive**

- Show system architecture diagram
- Open N8N workflow (visual appeal)
- Show Ollama running locally
- Demonstrate no internet connection

### **Minute 13-15: Cost & Scalability**

- Show hardware (\$0 procurement cost)
- Show power consumption (~21W total)
- Discuss scaling (add more T480s)
- Q&A

### **Demo Talking Points**

#### **For Judges:**

## **1. Security & Privacy**

- "Every single bit of data stays within this network"
- "No cloud providers, no third parties, complete data sovereignty"
- "Perfect for classified or sensitive government documents"

## **2. Cost Effectiveness**

- "Zero procurement cost - used existing hardware"
- "Under \$10/year electricity cost"
- "No per-query fees, no subscription costs"
- "Compare to commercial AI APIs: \$1000+/month for this volume"

## **3. Performance**

- "6-17 second end-to-end voice response"
- "Processing happens locally on this laptop"
- "Can handle 2-3 simultaneous users"
- "Scalable by adding more hardware"

## **4. Innovation**

- "Raspberry Pi as edge device (IoT/robotics showcase)"
- "Distributed architecture (real-world design pattern)"
- "Voice-enabled RAG (cutting-edge feature)"
- "Open-source stack (no vendor lock-in)"

## **5. Practical Applications**

- "Policy Q&A systems for government offices"
- "Internal knowledge bases"
- "Citizen service chatbots (offline capable)"
- "Emergency response systems (works without internet)"

## **Backup Plans**

**If something fails during demo:**

## **1. Voice doesn't work:**

- Fall back to text chat
- Explain: "Voice is experimental, text is rock-solid"

## **2. Network issues:**

- Have video recording of working system
- Show architecture diagrams and code

## **3. Slow response times:**

- Explain: "This is the small model, production would use dedicated GPU"
- Show performance benchmarks on slides

## **4. Questions you can't answer:**

- "Great question! This is a proof of concept, and that's exactly the kind of feedback we need for v2"
- 

# **Part 13: Post-Demo Evolution Path**

## **Short-term Upgrades (1-3 months)**

### **1. Add GPU to T480 (eGPU via Thunderbolt)**

- Cost: \$300-500 (used GTX 1070/1080)
- Performance: 5-10x faster inference
- Response time: 1-3 seconds (production-ready)

### **2. Scale horizontally**

- Add second T480 or similar laptop
- Load balancing via NGINX
- Support 10-20 concurrent users

### **3. Better models**

- Llama 3 70B (with GPU)
- Mixtral 8x7B
- Domain-specific fine-tuned models

## **Medium-term (3-6 months)**

## 1. Professional hardware

- Used server: Dell R730 (~\$500-800)
- 64-128GB RAM, Xeon CPUs
- Add NVIDIA Tesla P40 (\$200 used)
- Support 50+ users

## 2. Enhanced features

- Multi-language support
- Advanced RAG (HyDE, reranking)
- User authentication
- Conversation memory

## 3. Enterprise deployment

- Kubernetes cluster
  - High availability
  - Disaster recovery
  - Monitoring dashboard
- 

## Part 14: Complete Cost Breakdown

Your Setup (RPi 3B + T480)

**Hardware (Already Owned):**

- Raspberry Pi 3B: \$0
- Lenovo T480: \$0
- Gigabit switch: \$0 (or \$15 if needed)
- Ethernet cables: \$0 (or \$5)
- **Total Hardware: \$0-20**

**Software:**

- Everything is open-source: \$0

**Initial Setup Time:**

- RPi setup: 2 hours
- T480 setup: 8 hours
- Testing: 4 hours
- **Total: ~14 hours** (2 days part-time)

## **Operating Costs (Annual):**

- Electricity ( $21\text{W} \times 24\text{h} \times 365\text{d} \times \$0.12/\text{kWh}$ ): ~\$22
- SD card replacement (every 2-3 years): \$4/year
- Maintenance: \$0
- **Total OPEX: ~\$26/year**

## **5-Year Total Cost of Ownership:**

- CAPEX: \$0-20
- OPEX: \$130
- **Total: \$130-150** (vs \$60,000+ for commercial AI API services)

## **Dell Optiplex Comparison**

### **Hardware Cost:**

- Dell Optiplex 7050 Micro (i5-7500, 16GB): \$220
- Storage upgrade: \$0 (256GB included)
- Accessories: \$20
- **Total: \$240**

### **Performance:**

- 15-20% faster than T480
- Single device (simpler)
- Less impressive demo (no RPi showcase)

**Verdict: Not worth \$240 extra given your scenario**

---

## **Part 15: Final Recommendation & Action Plan**

**FINAL VERDICT: Use RPi 3B + T480 Setup**

**Why this is the right choice:**

1.  **Zero procurement cost** - You own the hardware
2.  **Impressive demo** - RPi visible, shows distributed architecture
3.  **Performance adequate** - 7-18s is acceptable for government demo
4.  **Voice functionality** - Fully local, no external APIs
5.  **Educational value** - Great for robotics/IoT fair
6.  **Scalable story** - Easy to explain growth path
7.  **Power efficient** - 21W total (green computing angle)
8.  **Real-world architecture** - Edge + backend pattern

### The Dell Optiplex would give you:

- Marginally faster (5-15s vs 7-18s) - **Not significant**
- Single device - **Less interesting demo**
- \$240 cost - **Not worth it for your scenario**

## 7-Day Implementation Plan

### Day 1: Hardware Setup

- Morning: Flash DietPi on RPi, configure network
- Afternoon: Install Arch Linux on T480, configure network
- Evening: Test connectivity between devices

### Day 2: Core Services

- Morning: Install Ollama, download models
- Afternoon: Install and test Whisper.cpp
- Evening: Install and test Piper TTS

### Day 3: Workflow Engine

- Morning: Install N8N
- Afternoon: Create document ingestion workflow
- Evening: Create RAG query workflow

### Day 4: Backend API

- Morning: Create FastAPI backend
- Afternoon: Integrate all services
- Evening: Test API endpoints

### Day 5: Frontend

- Morning: Deploy HTML/CSS/JS to RPi
- Afternoon: Configure NGINX, SSL
- Evening: Test complete flow

## Day 6: Testing & Optimization

- Morning: Performance benchmarks
- Afternoon: Load test, fix issues
- Evening: Document upload and query tests

## Day 7: Demo Preparation

- Morning: Create demo documents
- Afternoon: Rehearse demo script
- Evening: Final checks, backup plans

## Next Steps (Start Today!)

1. **Download ISOs:**
    - DietPi for RPi 3B
    - Arch Linux for T480
  2. **Prepare SD card for RPi**
  3. **Backup T480 before wiping**
  4. **Clone this document** for reference during installation
  5. **Set up test environment** on one device first
- 

## Part 16: Judges' FAQs (Be Prepared)

**Q: Why not use cloud services like ChatGPT API?** A: "This demo is for government applications where data privacy is paramount. Using cloud APIs means sending sensitive documents to third parties. Our system keeps 100% of data on-premises and works without internet."

**Q: Is 7-18 seconds response time too slow?** A: "For a proof of concept running on a \$0 laptop with no GPU, this is excellent. Production systems would add a GPU card (\$300) and reduce response time to 1-3 seconds. Compare that to \$1000+/month cloud costs."

**Q: Why use a Raspberry Pi if the T480 does all the work?** A: "The Pi demonstrates edge computing architecture - in production, we could have multiple Pis at different locations (offices, kiosks) all connecting to a central AI backend. It's a realistic distributed system design."

**Q: How secure is this really?** A: "Completely air-gapped - no external network connections possible. All data stays on local hardware. Perfect for classified documents. We can add encryption, authentication, and audit logs for production."

**Q: Can it scale to 100+ users?** A: "This demo supports 2-3 concurrent users. Production scaling is straightforward: add more T480-class machines (or better hardware) behind the same architecture. Each \$500 server can handle 20-30 users."

**Q: What about model accuracy?** A: "We're using Llama 3 8B, which is comparable to GPT-3.5 for most tasks. For better accuracy, we can use larger models (70B parameters) with GPU hardware. The RAG approach helps by grounding answers in your documents."

**Q: How do you update the AI models?** A: "Models are updated manually via USB drive in air-gapped mode, or automatically if we allow controlled internet access for updates only. We can schedule monthly updates, similar to antivirus definitions."

---

## Conclusion

You have an **excellent opportunity** here with zero procurement cost and impressive demo potential. The RPi + T480 setup gives you:

- Professional distributed architecture
- Full voice capabilities (local STT + TTS)
- Secure air-gapped operation
- Under \$30/year operating costs
- Perfect showcase for robotics/IoT/AI fair

**This is absolutely worth building and will impress government stakeholders!**

The 7-day implementation plan is realistic, and you'll have a working system that demonstrates cutting-edge concepts while staying completely secure and cost-effective.

Good luck with your demo! 